Do links to retracted papers predict future retractions? An out-of-sample test on the open biomedical citation graph

Hossein Mohasel Arjomandi*

Abstract

Tools that screen manuscripts for research-integrity problems flag papers that cite retracted work, and retrospective studies report that such papers are retracted more often. This study tests the signal prospectively on open data. Predictors were computed on the January 2024 OpenCitations graph using the retractions known at that date, frozen, and compared with retractions in the Retraction Watch database from February 2024 to October 2026. The design was pre-registered. Among 10.5 million biomedical papers published 2015–2023, 2.1% cited a paper already retracted, and 4,323 were retracted later (0.041%). Citing a retracted paper raised the later retraction rate from 0.035% to 0.27% with one such reference and 0.83% with three or four. Within journal and publication year the odds ratio was 2.99 (95% CI 2.73–3.27), so the signal is not explained by retraction waves in journals. As a screen it is weak: 0.9% of the 1,000 highest-ranked papers were retracted later, and 84% of later retractions cite no retracted paper. When the time point was moved into the late-2023 retraction wave, the journal's own retraction rate became the far better screen. NIH-funded papers cited already-retracted work as often as other papers in the same journal and year. Links to retracted papers are a useful flag but not a screen on their own.

Keywords: retraction, research integrity, citation networks, paper mills, OpenCitations, Retraction Watch

1 Introduction

Retractions have become far more frequent. In the Retraction Watch database, 13,210 papers have a first retraction notice dated 2023, more than in any earlier year, and papers retracted for paper-mill reasons rose from 246 in 2020 to 6,782 in 2023 [1, 2]. Because retractions are slow, often years after publication, editors and integrity teams look for early signals. One signal is already in use: the "Feet of Clay" detector of the Problematic Paper Screener flags articles whose reference lists include retracted papers [3]. Network studies support the idea. Retracted papers cluster in citation networks [4, 5], and in a retrospective analysis of more than 50,000 Web of Science retractions, articles citing retracted works had substantially higher subsequent retraction rates [6].

That association has not been tested in the way a screening tool is used: with information frozen at one moment and outcomes observed afterwards. A retrospective association can also arise from journals alone. Retractions arrive in waves within journals and special issues, and papers in the same venue cite one another, so a paper that cites a retracted paper may simply sit in a journal that is being cleaned up. Machine-learning studies of retraction prediction so far have used text and metadata, not citation links, and random rather than temporal splits [7].

*Microsoft, Redmond, WA, USA; hosseinm@microsoft.com. The views expressed in this article are the author's own and do not necessarily reflect those of Microsoft.

This article tests the signal prospectively on open data. All predictors are computed on the January 2024 OpenCitations graph [8] with the retractions known at that time, frozen with a hash, and then compared with retractions issued from February 2024 to October 2026, using the Retraction Watch database [1]. The scores, metrics and decision rule were pre-registered. The contributions are:

- a prospective, pre-registered test of whether citing (or being cited by) a retracted paper predicts later retraction, on 10.5 million biomedical papers;
- a separation of the neighbourhood signal from journal-level clustering, using strata of journal and publication year;
- an assessment of its value as a screen (precision and lift) against a journal baseline and a model fitted on an earlier window;
- supporting measurements of community concentration, NIH-funded exposure to retracted work and citations around retraction, all from open data, with code and run records.

2 Related work

Citation of retracted papers. Retracted papers continue to be cited after retraction, often without acknowledging it [9, 10]; see Teixeira da Silva et al. [11] for a review. Attention to retracted papers and its decline after retraction have been measured across platforms [12] and related to how widely known a finding was [13]. Large open resources now document post-retraction citations in biomedicine, including citation contexts [9, 14]. Case studies traced how far a single retracted result spreads through direct and indirect citations [15], and harm has been quantified by citation distance from retracted work [16].

Retractions in citation networks. In molecular biology, retracted papers form citation communities with high self-citation [4]. Paper-mill retractions reference and are cited by other later-retracted papers [5]. Across Web of Science, retractions cluster in citation networks, and citing retracted works is associated with higher subsequent retraction rates [6]. These studies are retrospective and do not separate the network signal from journal-level clustering.

Predicting retraction. Classifiers trained on titles, abstracts and metadata reach accuracies around 0.6–0.7 on balanced samples, with a random split [7]. The Problematic Paper Screener flags papers that cite retracted work for reassessment [3]; how well such flags anticipate later retractions has not been reported.

Funding. Misconduct-related retractions of NIH-funded papers carry direct costs [17], and retractions are followed by less funding and fewer citations in the affected fields [18]. How often NIH-funded papers themselves cite retracted work has not, to my knowledge, been measured.

Open citation data. OpenCitations provides open citation links and metadata [19–21]; its coverage differs from commercial indexes [22].

3 Data

Citation graph. Citations come from the parsed OpenCitations snapshot of January 2024 released through the Illinois Data Bank [8], built from the OpenCitations Index [19, 20]. It holds 1,962,840,983 citation links between 107,030,533 bibliographic resources. Identifiers (DOI, PMID) and journals were taken from the OpenCitations Meta dump [21, 23]: 26,582,951 graph nodes carry a PMID and 91,379,362 a DOI. Counting records and citing links by publication month shows that the snapshot is essentially complete up to September 2023: 240,000–300,000 records and 8.7–10.3 million citing links per month from January to September 2023, then 58,434 records and 2.4 million links in October and fewer afterwards.

Retractions. Retraction labels come from the Retraction Watch database distributed by Crossref [1]. One copy was frozen for the whole study (repository commit of 30 October 2026, 73,519 records). Only records whose nature is *Retraction* are used; expressions of concern and corrections are ignored. Records were grouped into papers by DOI, or by PMID when no DOI is given, keeping the earliest retraction notice. This gives 62,735 retracted papers. A paper is labelled *paper mill* when any of its Retraction Watch reasons contains "Paper Mill" (11,716 papers).

PubMed. The 2026 PubMed baseline (39,994,988 records) supplies PMID–DOI links, the PubMed retraction flags and grant agencies. A paper is *NIH-funded* when one of its grants lists an agency containing "NIH" (3,341,220 records). A graph node is *biomedical* when it has a PMID in OpenCitations Meta or its DOI appears in PubMed (31,464,435 nodes).

Matching. Retraction Watch papers were matched to graph nodes by DOI and by PMID. Of the 62,735 retracted papers, 52,254 (83.3%) match at least one node. Among Retraction Watch papers with a PMID, 88.8% are also flagged as retracted in PubMed; Retraction Watch is used as the single label source and PubMed as a check. In a random sample of 300 DOI-matched papers, the Retraction Watch and OpenCitations titles agreed for all 292 papers that had a title in OpenCitations Meta (differences were retraction prefixes such as "Notice of Retraction" or truncation).

4 Methods

4.1 Design

The test is prospective in the sense that matters for a screening tool: all predictors use only information available at a fixed time point T_0, and outcomes are retractions issued after T_0. The main time point is $T_0 = 31$ January 2024. Only 116,123 citing links in the graph (0.006%) carry a publication date after January 2024, so post-T_0 information can enter the predictors only negligibly. Because the graph is complete only to about September 2023, every analysis was repeated with $T_0' = 30$ September 2023. The design, scores, metrics and decision rule were pre-registered in a file whose SHA-256 hash was recorded before any score was computed; the analysis code checks the hashes of the frozen score files before reading any outcome.

Candidates and outcome. Candidates are biomedical nodes published 2015–2023 that have a journal, at least one reference in the graph, and no retraction notice dated on or before T_0 (10,495,850 papers). A candidate is *positive* if its first retraction notice falls between T_0 and 30 October 2026. Let R_0 denote the 44,401 graph nodes retracted on or before T_0.

Scores. For a candidate p with $d(p)$ references in the graph: $S_1(p)$ is the number of references of p that are in R_0; $S_2(p) = S_1(p)/d(p)$; $S_3(p)$ is the number of citations p received from papers in R_0; and $S_4(p)$ is the number of references of p to paper-mill retractions in R_0. The journal baseline is a smoothed retraction rate,

$$B(p) = \frac{r_j + m\,g}{n_j + m},$$

where j is the journal of p, n_j is the number of biomedical items in j published from 2010 to 2023, r_j is the number of those in R_0, g is the overall rate $\sum_j r_j / \sum_j n_j$, and $m = 100$.

4.2 Analyses

Association. Positive rates with Wilson intervals [24] are reported by S_1 level, and risk ratios for $S_1 \geq 1$, $S_3 \geq 1$ and $S_4 \geq 1$.

Within journal (primary test). Retractions cluster by journal and publication year, for example in mass retractions of special issues. To separate the neighbourhood signal from these clusters, the Mantel–Haenszel odds ratio [25] of exposure ($S_1 \geq 1$) on outcome was computed with strata defined by journal × publication year, with the Robins–Breslow–Greenland confidence interval [26]. The pre-registered decision rule was that the signal holds within journal if the 95% interval lies above 1.

Screening value. Candidates were ranked by (i) S_1, with ties broken by S_2 and then at random; (ii) the journal baseline B; and (iii) a logistic model with $\operatorname{logit} B$, $\log(1 + S_1)$, $\log(1 + S_3)$, an indicator of $S_1 \geq 1$ and indicators of age at the time point. The model was fitted only on an earlier window (predictors at 31 December 2021; outcomes up to T_0; candidates published 2013–2021) and applied unchanged. Precision at k, lift (precision divided by the base rate), recall and average precision [27] are reported with 95% intervals from 1,000 Poisson bootstrap resamples of the candidates.

Sensitivity analyses. The analyses were repeated (a) for candidates published 2015–2022, (b) without nodes whose DOI maps to more than one graph node, and (c) with T_0'.

Supporting analyses. Three descriptive analyses complement the test. (1) *Community concentration*: each retracted biomedical paper with a journal was matched to one non-retracted paper from the same journal and publication year, chosen by a fixed hash of its identifier; communities were found with the Leiden algorithm [28] (modularity, two iterations, seeds 1 and 2) on the subgraph formed by these papers and their neighbours, and the concentration of retracted papers over communities was compared with that of the controls under a permutation null that swaps labels within pairs. (2) *NIH exposure*: the share of biomedical papers published 2000–2023 that cite a paper in a year after its retraction, for NIH-funded and other papers, with a Mantel–Haenszel odds ratio stratified by journal and year. (3) *Citations around retraction*: citations per paper by year relative to the retraction year, for papers retracted 2000–2021 and matched controls.

Implementation. All joins and aggregations run in DuckDB [29] on one 8-core, 64 GB cloud machine. Every run writes its configuration, the hash of its script, its counts and checksums; per-paper scores are never released.

Table 1: Later retraction by number of references to papers retracted by T_0 (main analysis).

S_1	Candidates	Retracted later	Rate (95% CI)
0	10,277,289	3,640	0.035% (0.034–0.037)
1	187,735	502	0.27% (0.25–0.29)
2	23,200	119	0.51% (0.43–0.61)
3–4	6,175	51	0.83% (0.63–1.08)
≥ 5	1,451	11	0.76% (0.42–1.35)

5 Results

5.1 Exposure and outcomes

Of the 10,495,850 candidates, 218,561 (2.1%) cite at least one paper that had been retracted by T_0, and 172,149 are cited by one. Between February 2024 and October 2026, 4,323 candidates were retracted (0.041%), 418 of them for paper-mill reasons.

5.2 Association and dose

Candidates that cite a retracted paper were retracted later far more often (Table 1). The rate rises from 0.035% without a retracted reference to 0.27% with one, 0.51% with two and 0.83% with three or four; the 1,451 candidates with five or more retracted references show 0.76%, with a wide interval. The risk ratio for $S_1 \geq 1$ is 8.8 (95% CI 8.1–9.6). It is 17.8 (15.5–20.6) for citing a paper-mill retraction, and 6.8 (6.2–7.5) for being cited by a retracted paper.

5.3 Within journal and year

Part of this association reflects journals: retractions cluster in journals and publication years, and papers in the same journal cite one another. Within journal × publication year, the Mantel–Haenszel odds ratio for $S_1 \geq 1$ is 2.99 (95% CI 2.73–3.27; 1,201 informative strata holding 4,258 of the positives). By the pre-registered rule, the signal holds within journal. It holds for paper-mill outcomes (3.37, 2.65–4.29) and for other outcomes (2.92, 2.65–3.21). Being cited by a retracted paper gives a within-journal odds ratio of 3.76 (3.37–4.20). The estimate is unchanged without candidates whose DOI maps to several graph nodes (2.99) and slightly higher for papers published 2015–2022 (3.10, 2.82–3.41).

5.4 Screening value

Despite the strong relative effect, the signal is weak as a screen (Table 2). Most later retractions, 3,640 of 4,323 (84%), cite no retracted paper at all. Among the 1,000 candidates with the most retracted references, 9 were retracted later (0.9%, lift 22 over the base rate); among the top 10,000, 72 (0.72%), which is 1.7% of all positives. The journal rate alone ranks one positive among its top 1,000 (0.1%). The combined model, fitted on the earlier window, has the highest average precision (0.0037) but a lower precision at 1,000 than the S_1 ranking.

5.5 When the time point falls inside a retraction wave

With $T_0' = 30$ September 2023, the outcome window contains the mass retractions of late 2023 (7,578 positives, 3,141 of them paper-mill). The journal rate then becomes the best screen: 6.8%

Table 2: Screening performance in the main analysis (95% bootstrap intervals). Base rate 0.041%.

Ranking	Prec.@1,000	Lift@1,000	Prec.@10,000	Avg. precision
S_1	0.90% (0.40–1.50)	21.9	0.72% (0.56–0.89)	0.0013
Journal rate B	0.10% (0–0.30)	2.4	0.02%	0.0023
Combined model	0.40% (0.10–0.80)	9.7	0.79% (0.61–0.94)	0.0037

of its top 1,000 candidates were retracted (lift 94), against 1.1% for S_1. The within-journal odds ratio for $S_1 \geq 1$ falls to 2.00 (1.86–2.15), and to 1.27 (1.13–1.42) for paper-mill outcomes, while it stays at 2.97 for other outcomes. By January 2024 those waves had passed, and the journal rate placed only one future retraction among its top 1,000.

5.6 Supporting analyses

Community concentration. The 27,567 retracted biomedical papers with a journal, their controls from the same journal and year, and all their neighbours form a subgraph of 2.7 million papers and 57 million citations. Leiden found 197 and 200 communities with two seeds (modularity 0.65 and 0.66); the two partitions agree only moderately (adjusted Rand index 0.41), so both are reported. Retracted papers are more concentrated than their controls: the probability that two retracted papers share a community is 1.15 and 1.11 times that of two controls (permutation $p = 0.001$ for both, with a null mean of 1.00). The excess comes mainly from paper-mill retractions (1.68 and 1.52); for other retractions it is small (1.06 and 1.05). Retracted paper-mill papers thus cluster in the citation graph beyond what their journal and year imply, in line with Feng et al. [4] and Candal-Pedreira et al. [5].

NIH exposure. Of 2,332,727 NIH-funded biomedical papers published 2000–2023 in the graph, 10,872 (0.47%) cite at least one paper in a year after its retraction, compared with 0.48% of other biomedical papers. Within journal and year, NIH funding makes no practical difference (odds ratio 1.03, 1.00–1.05). NIH-funded papers cited 3,548 distinct retracted papers after their retraction.

Citations around retraction. For 11,788 papers retracted 2000–2021 and matched controls from the same journal and year, retracted papers received about 1.3 times the controls' citations in each of the five years before the retraction year, 0.66 times in the year after and 0.47 times two years after. Citations continue after retraction, at about two per paper in the second year, consistent with earlier studies [9, 10].

6 Discussion

The signal is real and not just the journal. Citing a paper that has already been retracted is followed by a markedly higher chance of retraction in the next three years, and the association survives comparison within journal and publication year (odds ratio about 3). This turns the retrospective association reported by Lou et al. [6] into a prospective one and shows that it is not explained by retraction waves within journals. The rate rises with the number of retracted references up to three or four, and citing a paper-mill retraction is the strongest single marker. Being cited by a retracted paper is at least as informative as citing one, which fits the picture of interlinked clusters of problematic papers [5].

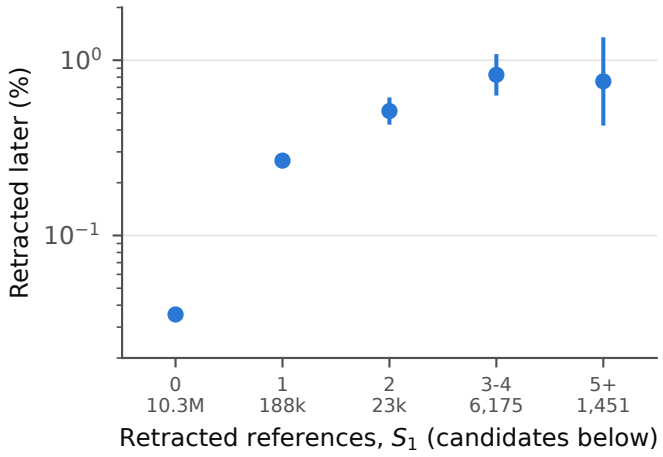


Figure 1: Share of candidates retracted between February 2024 and October 2026, by the number of their references to papers retracted by 31 January 2024 (95% Wilson intervals; log scale). Below each level: number of candidates.

But it is a weak screen on its own. A relative risk of nine on a base rate of 0.04% still means that almost all flagged papers are not retracted: about 1 in 110 of the top 1,000 by S_1. And 84% of later retractions cite no retracted paper at all, so the signal misses most cases. A "Feet of Clay" flag is therefore a reason to look, not evidence of a problem. Many papers cite retracted work for good reasons, such as reviews, critiques, or citations made before the retraction. Papers with five or more retracted references were not more likely to be retracted than papers with three or four; whether these are mostly reviews or commentaries that discuss retracted work was not examined.

Timing matters. The best screen depends on where the time point falls relative to retraction waves. Inside the late-2023 wave, the journal's own retraction rate identified future retractions with a lift of about 90, and the neighbourhood signal added little for paper-mill cases. Once the wave had passed, the journal rate carried almost no information about who would be retracted next, while the neighbourhood signal kept its strength. A practical screen should combine both and be refitted as waves come and go; the model fitted on outcomes up to January 2024, which put weight on the journal rate, ranked fewer later retractions at the very top than the reference count alone.

NIH-funded research. NIH-funded papers cite already-retracted work as often as other papers in the same journal and year. About 0.5% of biomedical papers published 2000–2023 do so: 103,349 papers in the graph, 10,872 of them NIH-funded. Whether these citations rely on the retracted findings cannot be read from citation links alone; citation-context resources [9, 14] could answer it.

Limitations. The test is limited to papers in the January 2024 OpenCitations snapshot, which depends on publishers depositing references with Crossref and is complete only to about September 2023; papers without deposited references cannot be scored, and 17% of retracted papers could not be matched to the graph. Retraction Watch is incomplete: 3,205 papers flagged as retracted in PubMed are missing from the copy used, three quarters of them with notices from 2024 onwards, so some later retractions are missing from the outcome labels. A retraction is not proof of misconduct; non-retracted candidates include papers that should have been retracted, so precision is a lower bound on the share of problematic papers among the flagged ones. The outcome window is under

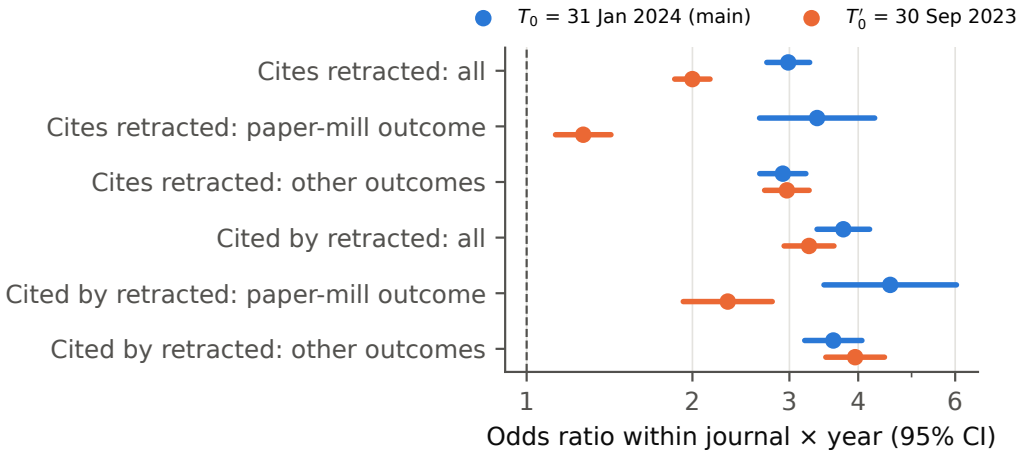


Figure 2: Mantel–Haenszel odds ratios within journal × publication year (95% Robins–Breslow–Greenland intervals), for citing a retracted paper ($S_1 \geq 1$) and being cited by one ($S_3 \geq 1$), with the main and the earlier time point.

three years, while the median time from publication to retraction for biomedical papers retracted after January 2024 was about 2.7 years, so many later retractions are not yet observed. NIH funding is taken from PubMed grant lists, which are incomplete. All results are associations.

7 Conclusion

On 10.5 million biomedical papers, citing a paper that had already been retracted predicted later retraction prospectively and within journal and year, with an odds ratio of about 3. The signal is a useful flag but a weak screen on its own: it ranks a few later retractions highly and misses most of them. Integrity screening should treat links to retracted work as one input, alongside journal-level signals that matter most during retraction waves. All data are open and the analysis can be rerun from the released code and records.

Data and code availability

Citation data: Mohasel Arjomandi et al. [8]; identifiers and journals: [23]; retractions: the Retraction Watch database distributed by Crossref (repository commit of 30 October 2026); PubMed 2026 baseline courtesy of the U.S. National Library of Medicine (it may not reflect the latest NLM data). Scripts, configuration, the pre-registration, run records with checksums, and the aggregate result tables: `https://github.com/HosseinMA96/retracted-neighbours-open-data`. The repository includes the pre-registration with its recorded hash, a step-by-step reproduction guide, and a map from every number in this article to the run that produced it. A derived dataset (retracted papers with their graph neighbourhoods, community identifiers and time-split labels; no per-paper scores for non-retracted papers, no author names): [TODO: Zenodo DOI].

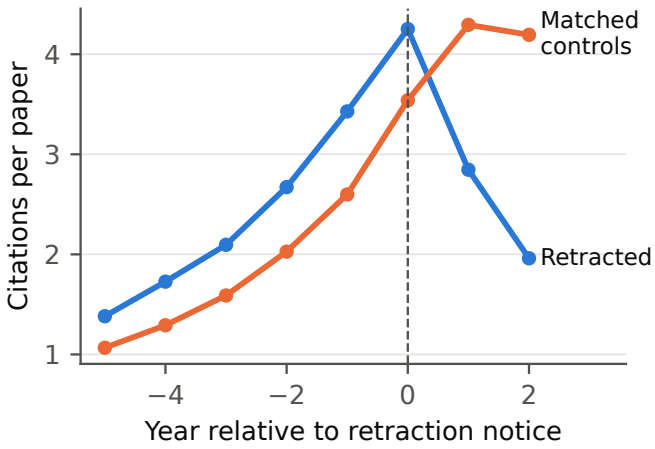


Figure 3: Citations per paper by year relative to the retraction notice, for 11,788 papers retracted 2000–2021 and controls matched on journal and publication year.

Declaration of generative AI use

Claude Opus 5.5 (Anthropic) was used as a coding and writing assistant: to write and run analysis scripts under the author's direction, to search the literature, to draft and edit the text, and to review drafts. All analyses, numbers and references were checked by the author against the run records and the original sources (DOIs), and the author takes full responsibility for the content.

References

[1] Martyn Rittman. Retraction watch retractions now in the crossref api. Crossref blog, 2025. URL `http://dx.doi.org/10.13003/692016`.

[2] Cristina Candal-Pedreira, Joseph S Ross, Alberto Ruano-Ravina, David S Egilman, Esteve Fernández, and Mónica Pérez-Ríos. Retracted papers originating from paper mills: cross sectional study. *BMJ*, 379:e071517, 2022. ISSN 1756-1833. doi: 10.1136/bmj-2022-071517. URL `http://dx.doi.org/10.1136/bmj-2022-071517`.

[3] Guillaume Cabanac, Cyril Labbé, and Alexander Magazinov. The 'problematic paper screener' automatically selects suspect publications for post-publication (re)assessment, 2022. URL `https://arxiv.org/abs/2210.04895`.

[4] Sida Feng, Lingzi Feng, Fang Han, Ye Zhang, Yanqing Ren, Lixue Wang, and Junpeng Yuan. Citation network analysis of retractions in molecular biology field. *Scientometrics*, 2024. ISSN 1588-2861. doi: 10.1007/s11192-024-05101-4. URL `http://dx.doi.org/10.1007/s11192-024-05101-4`.

[5] Cristina Candal-Pedreira, Carla Guerra-Tort, Alberto Ruano-Ravina, Fabián Freijedo-Farinas, Julia Rey-Brandariz, Joseph S. Ross, and Mónica Pérez-Ríos. Retracted papers originating from paper mills: a cross-sectional analysis of references and citations. *Journal of Clinical Epidemiology*, 172:111397, 2024. ISSN 0895-4356. doi: 10.1016/j.jclinepi.2024.111397. URL `http://dx.doi.org/10.1016/j.jclinepi.2024.111397`.

[6] Ying Lou, Zhengyi Zhou, and Menghui Li. Beyond retractions: A topic-based analysis of the impact and lifespan of retracted papers. *Journal of Data and Information Science*, 2026. ISSN 2543-683X. doi: 10.1515/jdis-2026-0064. URL http://dx.doi.org/10.1515/jdis-2026-0064.

[7] Aaron H. A. Fletcher and Mark Stevenson. Predicting retracted research: a dataset and machine learning approaches. *Research Integrity and Peer Review*, 10(1), 2025. ISSN 2058-8615. doi: 10.1186/s41073-025-00168-w. URL http://dx.doi.org/10.1186/s41073-025-00168-w.

[8] Hossein Mohasel Arjomandi, Dmitriy Korobskiy, and George Chacko. Parsed open citations and pubmed data, 2024. URL https://databank.illinois.edu/datasets/IDB-5216575.

[9] Tzu-Kun Hsiao and Jodi Schneider. Continued use of retracted papers: Temporal trends in citations and (lack of) awareness of retractions shown in citation contexts in biomedicine. *Quantitative Science Studies*, 2(4):1144–1169, 2021. ISSN 2641-3337. doi: 10.1162/qss_a_00155. URL http://dx.doi.org/10.1162/qss_a_00155.

[10] Jodi Schneider, Di Ye, Alison M. Hill, and Ashley S. Whitehorn. Continued post-retraction citation of a fraudulent clinical trial report, 11 years after it was retracted for falsifying data. *Scientometrics*, 125(3):2877–2913, 2020. ISSN 1588-2861. doi: 10.1007/s11192-020-03631-1. URL http://dx.doi.org/10.1007/s11192-020-03631-1.

[11] Jaime A. Teixeira da Silva, Karen Santos-d'Amorim, and Helmar Bornemann-Cimenti. The citation of retracted papers and impact on the integrity of the scientific biomedical literature. *Learned Publishing*, 38(2), 2025. ISSN 1741-4857. doi: 10.1002/leap.1667. URL http://dx.doi.org/10.1002/leap.1667.

[12] Hao Peng, Daniel M. Romero, and Emőke-Ágnes Horvát. Dynamics of cross-platform attention to retracted papers. *Proceedings of the National Academy of Sciences*, 119(25), 2022. ISSN 1091-6490. doi: 10.1073/pnas.2119086119. URL http://dx.doi.org/10.1073/pnas.2119086119.

[13] Shahan Ali Memon, Jevin D. West, and Cailin O'Connor. Are widely known findings easier to retract?, 2025. URL https://arxiv.org/abs/2504.15504.

[14] Stephan Joseph. Sanecite retraction census: post-retraction citations to 22,553 retracted biomedical papers (protocol, code and data), 2026. URL https://zenodo.org/doi/10.5281/zenodo.23766253.

[15] Paul E. van der Vet and Harm Nijveen. Propagation of errors in citation networks: a study involving the entire citation network of a widely cited paper published in, and later retracted from, the journal nature. *Research Integrity and Peer Review*, 1(1), 2016. ISSN 2058-8615. doi: 10.1186/s41073-016-0008-5. URL http://dx.doi.org/10.1186/s41073-016-0008-5.

[16] Yunyou Huang, Jiahui Zhao, Dandan Cui, Zhengxin Yang, Bingjie Xia, Qi Liang, Wenjing Liu, Li Ma, Suqin Tang, Tianyong Hao, Zhifei Zhang, Wanling Gao, and Jianfeng Zhan. Quantifying the dynamics of harm caused by retracted research, 2025. URL https://arxiv.org/abs/2501.00473.

[17] Andrew M Stern, Arturo Casadevall, R Grant Steen, and Ferric C Fang. Financial costs and personal consequences of research misconduct resulting in retracted publications. *eLife*, 3, 2014. ISSN 2050-084X. doi: 10.7554/elife.02956. URL http://dx.doi.org/10.7554/eLife.02956.

[18] Pierre Azoulay, Jeffrey L. Furman, Krieger, and Fiona Murray. Retractions. *Review of Economics and Statistics*, 97(5):1118–1136, 2015. ISSN 1530-9142. doi: 10.1162/rest_a_00469. URL http://dx.doi.org/10.1162/REST_a_00469.

[19] Silvio Peroni and David Shotton. Opencitations, an infrastructure organization for open scholarship. *Quantitative Science Studies*, 1(1):428–444, 2020. ISSN 2641-3337. doi: 10.1162/qss_a_00023. URL http://dx.doi.org/10.1162/qss_a_00023.

[20] Ivan Heibi, Silvio Peroni, and David Shotton. Software review: Coci, the opencitations index of crossref open doi-to-doi citations. *Scientometrics*, 121(2):1213–1228, 2019. ISSN 1588-2861. doi: 10.1007/s11192-019-03217-6. URL http://dx.doi.org/10.1007/s11192-019-03217-6.

[21] Arcangelo Massari, Fabio Mariani, Ivan Heibi, Silvio Peroni, and David Shotton. Opencitations meta. *Quantitative Science Studies*, 5(1):50–75, 2024. ISSN 2641-3337. doi: 10.1162/qss_a_00292. URL http://dx.doi.org/10.1162/qss_a_00292.

[22] Alberto Martín-Martín, Mike Thelwall, Enrique Orduna-Malea, and Emilio Delgado López-Cózar. Google scholar, microsoft academic, scopus, dimensions, web of science, and opencitations' coci: a multidisciplinary comparison of coverage via citations. *Scientometrics*, 126(1):871–906, 2020. ISSN 1588-2861. doi: 10.1007/s11192-020-03690-4. URL http://dx.doi.org/10.1007/s11192-020-03690-4.

[23] OpenCitations . Opencitations meta csv dataset of all bibliographic metadata, 2024. URL https://figshare.com/articles/dataset/OpenCitations_Meta_CSV_dataset_of_all_bibliographic_metadata/21747461/7.

[24] Edwin B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. ISSN 1537-274X. doi: 10.1080/01621459.1927.10502953. URL http://dx.doi.org/10.1080/01621459.1927.10502953.

[25] Nathan Mantel and William Haenszel. Statistical aspects of the analysis of data from retrospective studies of disease. *JNCI: Journal of the National Cancer Institute*, 22(4):719–748, 1959. doi: 10.1093/jnci/22.4.719.

[26] James Robins, Norman Breslow, and Sander Greenland. Estimators of the mantel-haenszel variance consistent in both sparse data and large-strata limiting models. *Biometrics*, 42(2):311, 1986. ISSN 0006-341X. doi: 10.2307/2531052. URL http://dx.doi.org/10.2307/2531052.

[27] Takaya Saito and Marc Rehmsmeier. The precision-recall plot is more informative than the roc plot when evaluating binary classifiers on imbalanced datasets. *PLOS ONE*, 10(3):e0118432, 2015. ISSN 1932-6203. doi: 10.1371/journal.pone.0118432. URL http://dx.doi.org/10.1371/journal.pone.0118432.

[28] V. A. Traag, L. Waltman, and N. J. van Eck. From louvain to leiden: guaranteeing well-connected communities. *Scientific Reports*, 9(1), 2019. ISSN 2045-2322. doi: 10.1038/s41598-019-41695-z. URL http://dx.doi.org/10.1038/s41598-019-41695-z.

[29] Mark Raasveldt and Hannes Mühleisen. Duckdb: an embeddable analytical database. In *Proceedings of the 2019 International Conference on Management of Data*, SIGMOD/PODS '19, page 1981–1984. ACM, 2019. doi: 10.1145/3299869.3320212. URL http://dx.doi.org/10.1145/3299869.3320212.